


Optimized memchr() Implementation For The Linux Kernel Up To ~4x Faster

Written by [Michael Larabel](#) in [Linux Kernel](#) on 10 July 2022 at 01:58 PM EDT. [25 Comments](#)



A set of proposed patches promise to make the Linux kernel's memchr() implementation faster for locating a character within a block of memory. In tests carried out by the developer, the new implementation can be nearly four times faster on large searches.

The latest memchr() patches sent out this Sunday for the Linux kernel aim to speed-up this function that can be around ~4x faster for long strings. Yu-Jen Chang who sent out the patches explained:

The original version of memchr() is implemented with the byte-wise comparing technique, which does not fully use 64-bits or 32-bits registers in CPU. We use word-wide comparing so that 8 characters can be compared at the same time on CPU. This code is base on David Laight's implementation.

We create two files to measure the performance. The first file contains on average 10 characters ahead the target character. The second file contains at least 1000 characters ahead the target character. Our implementation of "memchr()" is slightly better in the first test and nearly 4x faster than the original implementation in the second test.

This may be nice for micro-benchmarking memchr() but in the grand scheme probably won't equate to any material benefits for end-users especially as the biggest performance gains will be with very long strings. Checking the Linux kernel source tree today, there are around 129 times memchr() (or ~323 for different variations of memchr) is used ranging from just the Linux kernel tools to being used within file-system code and various drivers.

```
lib/test_hasan.c:      hasan_ptr_result = memchr(ptr, 'I', size + 1));
et/ceph/messenger.c:      delim_p = memchr(name, delim, namelen);
et/ceph/messenger.c:      colon_p = memchr(name, ':', namelen);
et/dns_resolver/dns_key.c:      opt = memchr(data, '#', datalen);
et/dns_resolver/dns_key.c:      next_opt = memchr(opt, '@', end - opt) ?: end;
et/dns_resolver/dns_key.c:      eq = memchr(opt, '=', opt_len);
et/mctp/device.c:      if (memchr(mdev->addr, addr->s_addr, mdev->num_addrs))
et/mctp/device.c:      pos = memchr(mdev->addr, addr->s_addr, mdev->num_addrs);
et/netlink/genetlink.c:      if (WARN_ON(memchr(grp->name, '\0', GENL_NAMSIZ) == NULL))
et/openvswitch/conntrack.c:      if (!memchr(helper, '\0', nla_len(a))) {
et/openvswitch/conntrack.c:      if (!memchr(info->timeout, '\0', nla_len(a))) {
et/rxrpc/rxkad.c:      q = memchr(p, 0, end - p);
et/tipc/netlink_compat.c:      return memchr(s, '\0', len) ? true : false;
cripts/asnl_compiler.c:      nl = memchr(line, '\n', end - buffer);
cripts/asnl_compiler.c:      while ((p = memchr(p, '-', nl - p)) {
cripts/asnl_compiler.c:      while ((q = memchr(q, '-', nl - q)) {
cripts/dtc/dtc-lexer.l:      fnstart = memchr(yytext, '\n', yyleng);
cripts/dtc/dtc-lexer.l:      if (memchr(fn.val, '\0', fn.len - 1))
cripts/dtc/libfdt/fdt_overlay.c:      fixup_end = memchr(value, '\0', len);
cripts/dtc/libfdt/fdt_overlay.c:      sep = memchr(fixup_str, ':', fixup_len);
cripts/dtc/libfdt/fdt_overlay.c:      sep = memchr(name, ':', fixup_len);
cripts/dtc/libfdt/fdt_overlay.c:      if (path_len < 1 || memchr(path, '\0', path_len) != &path[patl
))
cripts/dtc/libfdt/fdt_ro.c:      else if (!memchr(s, '@', len) && (p[len] == '@'))
cripts/dtc/libfdt/fdt_ro.c:      n = memchr(s, '\0', len);
cripts/dtc/libfdt/fdt_ro.c:      const char *q = memchr(path, '/', end - p);
cripts/dtc/libfdt/fdt_ro.c:      q = memchr(p, '/', end - p);
cripts/dtc/libfdt/fdt_ro.c:      p = memchr(strlist, '\0', listlen);
cripts/dtc/libfdt/libfdt_env.h:      const char *p = memchr(string, 0, max_count);
ecurity/tomoyo/util.c: const char *cp1 = memchr(filename, '/', len);
ecurity/tomoyo/util.c: const char *cp2 = memchr(filename, '.', len);
ols/lib/bpf/usdt.c:      name = (const char *)memchr(provider, '\0', data + len - provider);
ols/lib/bpf/usdt.c:      args = memchr(name, '\0', data + len - name);
ols/perf/builtin-trace.c:      sep = memchr(tok, '|', size);
ols/perf/util/color.c:      char *p = memchr(buf, '\n', count);
ols/perf/util/probe-event.c:      if (ptr && memchr(arg, '.', ptr - arg))
ols/perf/util/srcode.c:      while (p < end && (p = memchr(p, '\n', end - p)) != NULL) {
ols/perf/util/srcode.c:      while (p < end && (p = memchr(p, '\n', end - p)) != NULL) {
ols/perf/util/srcode.c:      p = memchr(l, '\n', sf->map + sf->maplen - 1);
ols/perf/util/symbol-elf.c:      name = (const char *)memchr(provider, '\0', data + len - provider);
ols/perf/util/symbol-elf.c:      args = memchr(name, '\0', data + len - name);
```

The patches for those interested can be found on the [Linux kernel mailing list](#) for review.

25 Comments

Related News

- [Fwupd 2.0.2 Allows Updating Firmware On Many More Devices](#)
- [Linux 6.13 Will Report The Number Of Hung Tasks Since Boot](#)
- [Sched_Ext Changes Merged For Linux 6.13 With LLC & NUMA Awareness](#)
- [Faster Raspberry Pi Graphics & Intel Xe3 Enablement Starts With Linux 6.13 DRM Changes](#)
- [Linux 6.13 "MM" Patches Bring Some Enticing Performance Optimizations](#)
- [Faster CRC32C & AEGIS-128 Crypto Performance On Linux 6.13 With Intel/AMD CPUs](#)

About The Author



Michael Larabel is the principal author of Phoronix.com and founded the site in 2004 with a focus on enriching the Linux hardware experience. Michael has written more than 20,000 articles covering the state of Linux hardware support, Linux performance, graphics drivers, and other topics. Michael is also the lead developer of the Phoronix Test Suite, Phormatic, and OpenBenchmarking.org automated benchmarking software. He can be followed via [Twitter](#), [LinkedIn](#), or contacted via [MichaelLarabel.com](#).

Popular News This Week

- [ReiserFS Has Been Deleted From The Linux Kernel](#)
- [Linux CoC Announces Decision Following Recent Bcachefs Drama](#)
- [Bcachefs Changes Rejected Reportedly Due To CoC, Kernel Future "Uncertain"](#)
- [Red Hat & Microsoft Bringing RHEL To WSL](#)
- [Linux 6.12 Released With Real-Time Capabilities, Sched_Ext, More AMD RDNA4 & More](#)
- [Linux 6.13 Will Report The Number Of Hung Tasks Since Boot](#)
- [9elements Takes Over Intel 1st Gen Xeon Scalable "Skylake" Support Within Coreboot](#)
- [Many Networking Changes In Linux 6.13 - One Line Of Code Helping WireGuard Performance](#)

Latest Linux News

- [Phoronix Premium Cyber Week Sale To Better Enjoy Our Linux Hardware Reviews & News](#)
- [Large Folio Patches For EXT4 Show Some Nice Performance Gains](#)
- [New AMD Zen 5 Perf Events Going Into Linux 6.13](#)
- [Raspberry Pi Pico 2 W Launches For \\$7](#)
- [Fwupd 2.0.2 Allows Updating Firmware On Many More Devices](#)
- [AMDXDNA Driver For Ryzen AI Now Ready To Appear In The Linux Kernel](#)
- [Virtual CPUFreq Driver Coming With Linux 6.13 For Better Power/Performance Within VMs](#)
- [Linux 6.13 Will Report The Number Of Hung Tasks Since Boot](#)
- [Microsoft Continues "Demikernel" Development LibOS For Kernel-Bypass I/O](#)
- [IBM Power11 CPUs Launching In 2025 - Linux 6.13 Preps KVM Nested Guests For Power11](#)

Latest Featured Articles

- [SilverStone XE360-SP5 & XE04-SP5 For Cooling AMD EPYC 9004/9005 4U Servers](#)
- [8 vs. 12 Channel DDR5-6000 Memory Performance With AMD 5th Gen EPYC](#)
- [New AMD ERAPS Feature Yields Additional Performance Gains On Zen 5](#)
- [Supermicro H13SSL-N For AMD EPYC 9005 "Turin" 1P Servers](#)
- [Apple M4 Mac Mini With macOS VS. Intel / AMD With Ubuntu Linux Performance](#)

Phoronix Media

Contact
Michael Larabel
OpenBenchmarking.org

Phoronix Premium

Support Phoronix
While Having Ad-Free Browsing,
Single-Page Article Viewing

Share

Facebook
Twitter

Legal Disclaimer, Privacy Policy, Cookies |
Contact

Copyright © 2004 - 2024 by Phoronix Media.

All trademarks used are properties of their respective owners. All rights reserved.